

■ Privacy-First AI Interviewer

Role: System Design
Date: 2026-08-01
Session ID: c42c1cc1-de6a-4001-92fb-e70ed3b616aa
Overall Score: N/A

Scorecard

Dimension	Score	Justification
Technical Depth	N/A	Scoring unavailable — model did not return valid JSON.
Communication Clarity	N/A	Scoring unavailable — model did not return valid JSON.
Confidence Fluency	N/A	Scoring unavailable — model did not return valid JSON.
Star Completeness	N/A	Scoring unavailable — model did not return valid JSON.
Problem Solving	N/A	Scoring unavailable — model did not return valid JSON.

■ ATS Resume Match Score: 1%

Matched Keywords: database, design, system
Missing Keywords: architect, balancing, caching, fault, indexing, latency
Suggestion: ATS score is low. Rewrite your resume to mirror the language used in the job description.

Full Transcript

Interviewer: I see you have significant experience building ML-powered systems, specifically with developing a real-time sentiment analysis dashboard deployed on AWS. To start, could you walk me through the end-to-end architecture of that project? I'm interested in understanding your initial design choices—how did you structure the pipeline from data ingestion all the way to the final user viewing the results? [TIME: You have 90 seconds to answer.]

Candidate: The architecture utilizes an Amazon Kinesis stream for real-time data ingestion which triggers an AWS Lambda function to analyze text using Amazon comprehend for instant sentiment scoring.

Interviewer: The pipeline sounds very efficient for instant scoring. But in a high-volume, real-time setting like this, what mechanism did you put in place to handle backpressure or ensure that if Comprehend throttled or failed temporarily, your Kinesis stream wouldn't lose messages and could resume processing cleanly? [TIME: You have 90 seconds to answer.]

Candidate: [skipped]

Interviewer: Given that we've discussed ingestion reliability, let's pivot to how the data ultimately becomes available for querying. Once the Lambda function successfully calculates the sentiment score, where does that structured result land? And critically, what trade-offs did you consider when choosing the final persistence layer—was it optimized more for write throughput or for low-latency dashboard reads, and why? [TIME: You have 90 seconds to answer.]

Candidate: [skipped]

Interviewer: Thinking about the entire lifecycle of this dashboard, we've covered real-time ingestion and persistence. Now, let's pivot to operations: When you needed to update or retrain

the ML model—how did your CI/CD pipeline handle deploying that new logic without downtime? And what specific monitoring metrics did you put in place to detect if the prediction quality or system throughput degraded over time? [TIME: You have 90 seconds to answer.]

Candidate: [skipped]

Interviewer: Thinking back over all the technical components—the Kinesis ingestion, Lambda processing, and the persistence layer—if you were asked to optimize this entire dashboard for maximum cost efficiency without sacrificing the core real-time functionality, what is the single biggest architectural change or trade-off you would make, and why? [TIME: You have 90 seconds to answer.]

Candidate: [skipped]

Summary & Next Steps

Scoring could not be completed automatically. Please review the transcript manually.